

Nonlinear Least-Squares Fitting

This chapter describes functions for multidimensional nonlinear least-squares fitting. There are generally two classes of **algorithms** for solving nonlinear least squares problems, which fall under line search methods and trust region methods. GSL currently implements only trust region methods and provides the user with full access to intermediate steps of the iteration. The user also has the ability to tune a number of parameters which affect low-level aspects of the algorithm which can help to accelerate convergence for the specific problem at hand. GSL provides two separate interfaces for nonlinear least squares fitting. The first is designed for small to moderate sized problems, and the second is designed for very large problems, which may or may not have significant sparse structure.

The header file `gsl_multifit_nlinear.h` contains prototypes for the multidimensional nonlinear fitting functions and related declarations relating to the small to moderate sized systems.

The header file `gsl_multilarge_nlinear.h` contains prototypes for the multidimensional nonlinear fitting functions and related declarations relating to large systems.

Overview

The problem of multidimensional nonlinear least-squares fitting requires the minimization of the squared residuals of n functions, f_i , in p parameters, x_i ,

$$\begin{aligned}\Phi(x) &= \frac{1}{2} \|f(x)\|^2 \\ &= \frac{1}{2} \sum_{i=1}^n f_i(x_1, \dots, x_p)^2\end{aligned}$$

In trust region methods, the objective (or cost) function $\Phi(x)$ is approximated by a model function $m_k(\delta)$ in the vicinity of some point x_k . The model function is often simply a second order Taylor series expansion around the point x_k , ie:

$$\Phi(x_k + \delta) \approx m_k(\delta) = \Phi(x_k) + g_k^T \delta + \frac{1}{2} \delta^T B_k \delta$$

where $g_k = \nabla \Phi(x_k) = J^T f$ is the gradient vector at the point x_k , $B_k = \nabla^2 \Phi(x_k)$ is the Hessian matrix at x_k , or some approximation to it, and J is the n -by- p Jacobian matrix

$$J_{ij} = \partial f_i / \partial x_j$$

In order to find the next step δ , we minimize the model function $m_k(\delta)$, but search for solutions only within a region where we trust that $m_k(\delta)$ is a good approximation to the objective function $\Phi(x_k + \delta)$. In other words, we seek a solution of the trust region subproblem (TRS)

$$\min_{\delta \in \mathbb{R}^p} m_k(\delta) = \Phi(x_k) + g_k^T \delta + \frac{1}{2} \delta^T B_k \delta, \quad \text{s.t.} \quad \|D_k \delta\| \leq \Delta_k$$

where $\Delta_k > 0$ is the trust region radius and D_k is a scaling matrix. If $D_k = I$, then the trust region is a ball of radius Δ_k centered at x_k . In some applications, the parameter vector x may have widely different scales. For example, one parameter might be a temperature on the order of 10^3 K, while another might be a length on the order of 10^{-6} m. In such cases, a spherical trust region may not be the best choice, since if Φ changes rapidly along directions with one scale, and more slowly along directions with a different scale, the model function m_k may be a poor approximation to Φ along the rapidly changing directions. In such problems, it may be best to use an elliptical trust region, by setting D_k to a diagonal matrix whose entries are designed so that the scaled step $D_k \delta$ has entries of approximately the same order of magnitude.

The trust region subproblem above normally amounts to solving a linear least squares system (or multiple systems) for the step δ . Once δ is computed, it is checked whether or not it reduces the objective function $\Phi(x)$. A useful statistic for this is to look at the ratio

$$\rho_k = \frac{\Phi(x_k) - \Phi(x_k + \delta_k)}{m_k(0) - m_k(\delta_k)}$$

where the numerator is the actual reduction of the objective function due to the step δ_k , and the denominator is the predicted reduction due to the model m_k . If ρ_k is negative, it means that the step δ_k increased the objective function and so it is rejected. If ρ_k is positive, then we have found a step which reduced the objective function and it is accepted. Furthermore, if ρ_k is close to 1, then this indicates that the model function is a good approximation to the objective function in the trust region, and so on the next iteration the trust region is enlarged in order to take more ambitious steps. When a step is rejected, the trust region is made smaller and the TRS is solved again. An outline for the general trust region method used by GSL can now be given.

Trust Region Algorithm

1. Initialize: given x_0 , construct $m_0(\delta)$, D_0 and $\Delta_0 > 0$
2. For $k = 0, 1, 2, \dots$
 - a. If converged, then stop
 - b. Solve TRS for trial step δ_k
 - c. Evaluate trial step by computing ρ_k
 - 1). if step is accepted, set $x_{k+1} = x_k + \delta_k$ and increase radius,
$$\Delta_{k+1} = \alpha \Delta_k$$
 - 2). if step is rejected, set $x_{k+1} = x_k$ and decrease radius,
$$\Delta_{k+1} = \frac{\Delta_k}{\beta}, \text{ goto 2(b)}$$
 - d. Construct $m_{k+1}(\delta)$ and D_{k+1}

GSL offers the user a number of different **algorithms** for solving the trust region subproblem in 2(b), as well as different choices of scaling matrices D_k and different methods of updating the trust region radius Δ_k . Therefore, while reasonable default methods are provided, the user has a lot of control to fine-tune the various steps of the algorithm for their specific problem.

Solving the Trust Region Subproblem (TRS)

Below we describe the methods available for solving the trust region subproblem. The methods available provide either exact or approximate solutions to the trust region subproblem. In all **algorithms** below, the Hessian matrix B_k is approximated as $B_k \approx J_k^T J_k$, where $J_k = J(x_k)$. In all methods, the solution of the TRS involves solving a linear least squares system involving the Jacobian matrix. For small to moderate sized problems (`gsl_multifit_nlinear` interface), this is accomplished by factoring the full Jacobian matrix, which is provided by the user, with the Cholesky, QR, or SVD decompositions. For large systems (`gsl_multilarge_nlinear` interface), the

user has two choices. One is to solve the system iteratively, without needing to store the full Jacobian matrix in memory. With this method, the user must provide a routine to calculate the matrix-vector products $J\mathbf{u}$ or $J^T\mathbf{u}$ for a given vector \mathbf{u} . This iterative method is particularly useful for systems where the Jacobian has sparse structure, since forming matrix-vector products can be done cheaply. The second option for large systems involves forming the normal equations matrix $J^T J$ and then factoring it using a Cholesky decomposition. The normal equations matrix is p -by- p , typically much smaller than the full n -by- p Jacobian, and can usually be stored in memory even if the full Jacobian matrix cannot. This option is useful for large, dense systems, or if the iterative method has difficulty converging.

Levenberg-Marquardt

There is a theorem which states that if δ_k is a solution to the trust region subproblem given above, then there exists $\mu_k \geq 0$ such that

$$(B_k + \mu_k D_k^T D_k) \delta_k = -g_k$$

with $\mu_k(\Delta_k - \|D_k \delta_k\|) = 0$. This forms the basis of the **Levenberg-Marquardt algorithm**, which controls the trust region size by adjusting the parameter μ_k rather than the radius Δ_k directly. For each radius Δ_k , there is a unique parameter μ_k which solves the TRS, and they have an inverse relationship, so that large values of μ_k correspond to smaller trust regions, while small values of μ_k correspond to larger trust regions.

With the approximation $B_k \approx J_k^T J_k$, on each iteration, in order to calculate the step δ_k , the following linear least squares problem is solved:

$$\begin{bmatrix} J_k \\ \sqrt{\mu_k} D_k \end{bmatrix} \delta_k = - \begin{bmatrix} f_k \\ 0 \end{bmatrix}$$

If the step δ_k is accepted, then μ_k is decreased on the next iteration in order to take a larger step, otherwise it is increased to take a smaller step. The **Levenberg-Marquardt algorithm** provides an exact solution of the trust region subproblem, but typically has a higher computational cost per iteration than the approximate methods discussed below, since it may need to solve the least squares system above several times for different values of μ_k .

Levenberg-Marquardt with Geodesic Acceleration

This method applies a so-called geodesic acceleration correction to the standard **Levenberg-Marquardt** step δ_k (Transtrum et al, 2011). By interpreting δ_k as a first order step along a geodesic in the model parameter space (ie: a velocity $\delta_k = v_k$), the geodesic acceleration a_k is a second order correction along the geodesic which is determined by solving the linear least squares system

$$\begin{bmatrix} J_k \\ \sqrt{\mu_k} D_k \end{bmatrix} a_k = - \begin{bmatrix} f_{vv}(x_k) \\ 0 \end{bmatrix}$$

where f_{vv} is the second directional derivative of the residual vector in the velocity direction v , $f_{vv}(x) = D_v^2 f = \sum_{\alpha\beta} v_\alpha v_\beta \partial_\alpha \partial_\beta f(x)$, where α and β are summed over the p parameters. The new total step is then $\delta'_k = v_k + \frac{1}{2}a_k$. The second order correction a_k can be calculated with a modest additional cost, and has been shown to dramatically reduce the number of iterations (and expensive Jacobian evaluations) required to reach convergence on a variety of different problems. In order to utilize the geodesic acceleration, the user must supply a function which provides the second directional derivative vector $f_{vv}(x)$, or alternatively the library can use a finite difference method to estimate this vector with one additional function evaluation of $f(x + hv)$ where h is a tunable step size (see the `h_fvv` parameter description).

Dogleg

This is Powell's dogleg method, which finds an approximate solution to the trust region subproblem, by restricting its search to a piecewise linear "dogleg" path, composed of the origin, the Cauchy point which represents the model minimizer along the steepest descent direction, and the Gauss-Newton point, which is the overall minimizer of the unconstrained model. The Gauss-Newton step is calculated by solving

$$J_k \delta_{gn} = -f_k$$

which is the main computational task for each iteration, but only needs to be performed once per iteration. If the Gauss-Newton point is inside the trust region, it is selected as the step. If it is outside, the method then calculates the Cauchy point, which is located along the gradient direction. If the Cauchy point is also outside the trust region, the method assumes that it is still far from the minimum and so proceeds along the gradient direction, truncating the step at the trust region boundary. If the Cauchy point is inside the trust region, with the Gauss-Newton point outside, the method uses a dogleg step, which is a linear combination of the gradient direction and the Gauss-Newton direction, stopping at the trust region boundary.

Double Dogleg

This method is an improvement over the classical dogleg **algorithm**, which attempts to include information about the Gauss-Newton step while the iteration is still far from the minimum. When the Cauchy point is inside the trust region and the Gauss-Newton point is outside, the method computes a scaled Gauss-Newton point and then takes a dogleg step between the Cauchy point and the scaled Gauss-Newton point. The scaling is calculated to ensure that the reduction in the model m_k is about the same as the reduction provided by the Cauchy point.

Two Dimensional Subspace

The dogleg methods restrict the search for the TRS solution to a 1D curve defined by the Cauchy and Gauss-Newton points. An improvement to this is to search for a solution using the full two dimensional subspace spanned by the Cauchy and Gauss-Newton directions. The dogleg path is of course inside this subspace, and so this method solves the TRS at least as accurately as the dogleg methods. Since this method searches a larger subspace for a solution, it can converge more quickly than dogleg on some problems. Because the subspace is only two dimensional, this method is very efficient and the main computation per iteration is to determine the Gauss-Newton point.

Steihaug-Toint Conjugate Gradient

One difficulty of the dogleg methods is calculating the Gauss-Newton step when the Jacobian matrix is singular. The Steihaug-Toint method also computes a generalized dogleg step, but avoids solving for the Gauss-Newton step directly, instead using an iterative conjugate gradient **algorithm**. This method performs well at points where the Jacobian is singular, and is also suitable for large-scale problems where factoring the Jacobian matrix could be prohibitively expensive.

Weighted Nonlinear Least-Squares

Weighted nonlinear least-squares fitting minimizes the function

$$\begin{aligned}\Phi(x) &= \frac{1}{2} \|f\|_W^2 \\ &= \frac{1}{2} \sum_{i=1}^n w_i f_i(x_1, \dots, x_p)^2\end{aligned}$$

where $W = \text{diag}(w_1, w_2, \dots, w_n)$ is the weighting matrix, and $\|f\|_W^2 = f^T W f$. The weights w_i are commonly defined as $w_i = 1/\sigma_i^2$, where σ_i is the error in the i -th measurement. A simple change of variables $\tilde{f} = W^{\frac{1}{2}} f$ yields $\Phi(x) = \frac{1}{2} \|\tilde{f}\|^2$, which is in the same form as the unweighted case. The user can either perform this transform directly on their function residuals and Jacobian, or use the `gsl_multifit_nlinear_winit()` interface which automatically performs the correct scaling. To manually perform this transformation, the residuals and Jacobian should be modified according to

$$\tilde{f}_i = \sqrt{w_i} f_i = \frac{f_i}{\sigma_i}$$

$$\tilde{J}_{ij} = \sqrt{w_i} \frac{\partial f_i}{\partial x_j} = \frac{1}{\sigma_i} \frac{\partial f_i}{\partial x_j}$$

For large systems, the user must perform their own weighting.

Tunable Parameters

The user can tune nearly all aspects of the iteration at allocation time. For the

`gsl_multifit_nlinear` interface, the user may modify the `gsl_multifit_nlinear_parameters` structure, which is defined as follows:

type `gsl_multifit_nlinear_parameters`

```
typedef struct
{
    const gsl_multifit_nlinear_trs *trs;          /* trust region subproblem method */
    const gsl_multifit_nlinear_scale *scale;      /* scaling method */
    const gsl_multifit_nlinear_solver *solver;    /* solver method */
    gsl_multifit_nlinear_fdtype fdtype;          /* finite difference method */
    double factor_up;                             /* factor for increasing trust radius */
    double factor_down;                           /* factor for decreasing trust radius */
    double avmax;                                 /* max allowed |a|/|v| */
    double h_df;                                 /* step size for finite difference Jacobian */
    /*
    double h_fvv;                               /* step size for finite difference fvv */
} gsl_multifit_nlinear_parameters;
```

For the `gsl_multilarge_nlinear` interface, the user may modify the

`gsl_multilarge_nlinear_parameters` structure, which is defined as follows:

type `gsl_multilarge_nlinear_parameters`

```
typedef struct
{
    const gsl_multilarge_nlinear_trs *trs;          /* trust region subproblem method */
    const gsl_multilarge_nlinear_scale *scale;      /* scaling method */
    const gsl_multilarge_nlinear_solver *solver;    /* solver method */
    gsl_multilarge_nlinear_fdtype fdtype;          /* finite difference method */
    double factor_up;                             /* factor for increasing trust radius */
    double factor_down;                           /* factor for decreasing trust radius */
    double avmax;                                 /* max allowed |a|/|v| */
    double h_df;                                 /* step size for finite difference
    Jacobian */
    double h_fvv;                               /* step size for finite difference fvv */
    size_t max_iter;                             /* maximum iterations for trs method */
    double tol;                                  /* tolerance for solving trs */
} gsl_multilarge_nlinear_parameters;
```

Each of these parameters is discussed in further detail below.

type `gsl_multifit_nlinear_trs`

type gsl_multilarge_nlinear_trs

The parameter `trs` determines the method used to solve the trust region subproblem, and may be selected from the following choices,

`gsl_multifit_nlinear_trs *gsl_multifit_nlinear_trs_lm`

`gsl_multilarge_nlinear_trs *gsl_multilarge_nlinear_trs_lm`

This selects the **Levenberg-Marquardt algorithm**.

`gsl_multifit_nlinear_trs *gsl_multifit_nlinear_trs_lmaccel`

`gsl_multilarge_nlinear_trs *gsl_multilarge_nlinear_trs_lmaccel`

This selects the **Levenberg-Marquardt algorithm** with geodesic acceleration.

`gsl_multifit_nlinear_trs *gsl_multifit_nlinear_trs_dogleg`

`gsl_multilarge_nlinear_trs *gsl_multilarge_nlinear_trs_dogleg`

This selects the **dogleg algorithm**.

`gsl_multifit_nlinear_trs *gsl_multifit_nlinear_trs_ddogleg`

`gsl_multilarge_nlinear_trs *gsl_multilarge_nlinear_trs_ddogleg`

This selects the **double dogleg algorithm**.

`gsl_multifit_nlinear_trs *gsl_multifit_nlinear_trs_subspace2D`

`gsl_multilarge_nlinear_trs *gsl_multilarge_nlinear_trs_subspace2D`

This selects the **2D subspace algorithm**.

`gsl_multilarge_nlinear_trs *gsl_multilarge_nlinear_trs_cgst`

This selects the **Steihaug-Toint conjugate gradient algorithm**. This method is available only for large systems.

type gsl_multifit_nlinear_scale

type gsl_multilarge_nlinear_scale

The parameter `scale` determines the diagonal scaling matrix D and may be selected from the following choices,

`gsl_multifit_nlinear_scale *gsl_multifit_nlinear_scale_more`

`gsl_multilarge_nlinear_scale *gsl_multilarge_nlinear_scale_more`

This damping strategy was suggested by Moré, and corresponds to $D^T D = \max(\text{diag}(J^T J))$, in other words the maximum elements of $\text{diag}(J^T J)$ encountered thus far in the iteration. This choice of D makes the problem scale-invariant, so that if the model parameters x_i are each scaled by an arbitrary constant, $\tilde{x}_i = a_i x_i$, then the sequence of iterates produced by the **algorithm** would be unchanged. This method can work very well in cases where the model parameters have widely different scales (ie: if some parameters are measured in nanometers, while others are measured in degrees Kelvin). This strategy has been proven effective on a large class of problems and so it is the library default, but it may not be the best choice for all problems.

`gsl_multifit_nlinear_scale *gsl_multifit_nlinear_scale_levenberg`

`gsl_multilarge_nlinear_scale *gsl_multilarge_nlinear_scale_levenberg`

This damping strategy was originally suggested by **Levenberg**, and corresponds to $D^T D = I$. This method has also proven effective on a large class of problems, but is not scale-invariant. However, some authors (e.g. Transtrum and Sethna 2012) argue that this choice is better for problems which are susceptible to parameter evaporation (ie: parameters go to infinity)

`gsl_multifit_nlinear_scale *gsl_multifit_nlinear_scale_marquardt`

`gsl_multilarge_nlinear_scale *gsl_multilarge_nlinear_scale_marquardt`

This damping strategy was suggested by **Marquardt**, and corresponds to $D^T D = \text{diag}(J^T J)$. This method is scale-invariant, but it is generally considered inferior to both the **Levenberg** and Moré strategies, though may work well on certain classes of problems.

`type gsl_multifit_nlinear_solver`

`type gsl_multilarge_nlinear_solver`

Solving the trust region subproblem on each iteration almost always requires the solution of the following linear least squares system

$$\begin{bmatrix} J \\ \sqrt{\mu} D \end{bmatrix} \delta = - \begin{bmatrix} f \\ 0 \end{bmatrix}$$

The `solver` parameter determines how the system is solved and can be selected from the following choices:

`gsl_multifit_nlinear_solver *gsl_multifit_nlinear_solver_qr`

This method solves the system using a rank revealing QR decomposition of the Jacobian J . This method will produce reliable solutions in cases where the Jacobian is rank deficient or near-singular but does require about twice as many operations as the Cholesky method discussed below.

`gsl_multifit_nlinear_solver *gsl_multifit_nlinear_solver_cholesky`

`gsl_multilarge_nlinear_solver *gsl_multilarge_nlinear_solver_cholesky`

This method solves the alternate normal equations problem

$$(J^T J + \mu D^T D) \delta = -J^T f$$

by using a Cholesky decomposition of the matrix $J^T J + \mu D^T D$. This method is faster than the QR approach, however it is susceptible to numerical instabilities if the Jacobian matrix is rank deficient or near-singular. In these cases, an attempt is made to reduce the condition number of the matrix using Jacobi preconditioning, but for highly ill-conditioned problems the QR approach is better. If it is known that the Jacobian matrix is well conditioned, this method is accurate and will perform faster than the QR approach.

`gsl_multifit_nlinear_solver *gsl_multifit_nlinear_solver_mcholesky`

`gsl_multilarge_nlinear_solver *gsl_multilarge_nlinear_solver_mcholesky`

This method solves the alternate normal equations problem

$$(J^T J + \mu D^T D) \delta = -J^T f$$

by using a modified Cholesky decomposition of the matrix $J^T J + \mu D^T D$. This is more suitable for the dogleg methods where the parameter $\mu = 0$, and the matrix $J^T J$ may be ill-conditioned or indefinite causing the standard Cholesky decomposition to fail. This method is based on Level 2 BLAS and is thus slower than the standard Cholesky decomposition, which is based on Level 3 BLAS.

`gsl_multifit_nlinear_solver *gsl_multifit_nlinear_solver_svd`

This method solves the system using a singular value decomposition of the Jacobian J . This method will produce the most reliable solutions for ill-conditioned Jacobians but is also the slowest solver method.

`type gsl_multifit_nlinear_fdtype`

The parameter `fdtype` specifies whether to use forward or centered differences when approximating the Jacobian. This is only used when an analytic Jacobian is not provided to the solver. This parameter may be set to one of the following choices.

`GSL_MULTIFIT_NLINEAR_FWDIFF`

This specifies a forward finite difference to approximate the Jacobian matrix. The Jacobian matrix will be calculated as

$$J_{ij} = \frac{1}{\Delta_j} (f_i(x + \Delta_j e_j) - f_i(x))$$

where $\Delta_j = h|x_j|$ and e_j is the standard j -th Cartesian unit basis vector so that $x + \Delta_j e_j$ represents a small (forward) perturbation of the j -th parameter by an amount Δ_j . The perturbation Δ_j is proportional to the current value $|x_j|$ which helps to calculate an accurate Jacobian when the various parameters have different scale sizes. The value of h is specified by the `h_df` parameter. The accuracy of this method is $O(h)$, and evaluating this matrix requires an additional p function evaluations.

`GSL_MULTIFIT_NLINEAR_CTRDIFF`

This specifies a centered finite difference to approximate the Jacobian matrix. The Jacobian matrix will be calculated as

$$J_{ij} = \frac{1}{\Delta_j} \left(f_i(x + \frac{1}{2}\Delta_j e_j) - f_i(x - \frac{1}{2}\Delta_j e_j) \right)$$

See above for a description of Δ_j . The accuracy of this method is $O(h^2)$, but evaluating this matrix requires an additional $2p$ function evaluations.

`double factor_up`

When a step is accepted, the trust region radius will be increased by this factor. The default value is 3.

`double factor_down`

When a step is rejected, the trust region radius will be decreased by this factor. The default value is 2.

`double avmax`

When using geodesic acceleration to solve a nonlinear least squares problem, an important parameter to monitor is the ratio of the acceleration term to the velocity term,

$$\frac{\|a\|}{\|v\|}$$

If this ratio is small, it means the acceleration correction is contributing very little to the step. This could be because the problem is not “nonlinear” enough to benefit from the acceleration. If the ratio is large (> 1) it means that the acceleration is larger than the velocity, which shouldn't happen since the step represents a truncated series and so the second order term a should be smaller than

the first order term v to guarantee convergence. Therefore any steps with a ratio larger than the parameter `avmax` are rejected. `avmax` is set to 0.75 by default. For problems which experience difficulty converging, this threshold could be lowered.

`double h_df`

This parameter specifies the step size for approximating the Jacobian matrix with finite differences. It is set to $\sqrt{\epsilon}$ by default, where ϵ is `GSL_DBL_EPSILON`.

`double h_fvv`

When using geodesic acceleration, the user must either supply a function to calculate $f_{vv}(x)$ or the library can estimate this second directional derivative using a finite difference method. When using finite differences, the library must calculate $f(x + hv)$ where h represents a small step in the velocity direction. The parameter `h_fvv` defines this step size and is set to 0.02 by default.

Initializing the Solver

`type gsl_multifit_nlinear_type`

This structure specifies the type of **algorithm** which will be used to solve a nonlinear least squares problem. It may be selected from the following choices,

`gsl_multifit_nlinear_type *gsl_multifit_nlinear_trust`

This specifies a trust region method. It is currently the only implemented nonlinear least squares method.

`gsl_multifit_nlinear_workspace *gsl_multifit_nlinear_alloc(const gsl_multifit_nlinear_type *T, const gsl_multifit_nlinear_parameters *params, const size_t n, const size_t p)`

`gsl_multilarge_nlinear_workspace *gsl_multilarge_nlinear_alloc(const gsl_multilarge_nlinear_type *T, const gsl_multilarge_nlinear_parameters *params, const size_t n, const size_t p)`

These functions return a pointer to a newly allocated instance of a derivative solver of type `T` for `n` observations and `p` parameters. The `params` input specifies a tunable set of parameters which will affect important details in each iteration of the trust region subproblem **algorithm**. It is recommended to start with the suggested default parameters (see `gsl_multifit_nlinear_default_parameters()` and `gsl_multilarge_nlinear_default_parameters()`) and then tune the parameters once the code is working correctly. See [Tunable Parameters](#) for descriptions of the various parameters. For example, the following code creates an instance of a **Levenberg-Marquardt** solver for 100 data points and 3 parameters, using suggested defaults:

```
const gsl_multifit_nlinear_type * T = gsl_multifit_nlinear_trust;
gsl_multifit_nlinear_parameters params = gsl_multifit_nlinear_default_parameters();
gsl_multifit_nlinear_workspace * w = gsl_multifit_nlinear_alloc (T, &params, 100, 3);
```

The number of observations `n` must be greater than or equal to parameters `p`.

If there is insufficient memory to create the solver then the function returns a null pointer and the error handler is invoked with an error code of `GSL_ENOMEM`.

`gsl_multifit_nlinear_parameters gsl_multifit_nlinear_default_parameters(void)`

`gsl_multilarge_nlinear_parameters gsl_multilarge_nlinear_default_parameters(void)`

These functions return a set of recommended default parameters for use in solving nonlinear least squares problems. The user can tune each parameter to improve the performance on their particular problem, see [Tunable Parameters](#).

```
int gsl_multifit_nlinear_init(const gsl_vector *x, gsl_multifit_nlinear_fdf *fdf, gsl_multifit_nlinear_workspace *w)
```

```
int gsl_multifit_nlinear_winit(const gsl_vector *x, const gsl_vector *wts, gsl_multifit_nlinear_fdf *fdf, gsl_multifit_nlinear_workspace *w)
```

```
int gsl_multilarge_nlinear_init(const gsl_vector *x, gsl_multilarge_nlinear_fdf *fdf, gsl_multilarge_nlinear_workspace *w)
```

These functions initialize, or reinitialize, an existing workspace `w` to use the system `fdf` and the initial guess `x`. See [Providing the Function to be Minimized](#) for a description of the `fdf` structure.

Optionally, a weight vector `wts` can be given to perform a weighted nonlinear regression. Here, the weighting matrix is $W = \text{diag}(w_1, w_2, \dots, w_n)$.

```
void gsl_multifit_nlinear_free(gsl_multifit_nlinear_workspace *w)
```

```
void gsl_multilarge_nlinear_free(gsl_multilarge_nlinear_workspace *w)
```

These functions free all the memory associated with the workspace `w`.

```
const char *gsl_multifit_nlinear_name(const gsl_multifit_nlinear_workspace *w)
```

```
const char *gsl_multilarge_nlinear_name(const gsl_multilarge_nlinear_workspace *w)
```

These functions return a pointer to the name of the solver. For example:

```
printf ("w is a '%s' solver\n", gsl_multifit_nlinear_name (w));
```

would print something like `w is a 'trust-region' solver`.

```
const char *gsl_multifit_nlinear_trs_name(const gsl_multifit_nlinear_workspace *w)
```

```
const char *gsl_multilarge_nlinear_trs_name(const gsl_multilarge_nlinear_workspace *w)
```

These functions return a pointer to the name of the trust region subproblem method. For example:

```
printf ("w is a '%s' solver\n", gsl_multifit_nlinear_trs_name (w));
```

would print something like `w is a 'levenberg-marquardt' solver`.

Providing the Function to be Minimized

The user must provide n functions of p variables for the minimization **algorithm** to operate on. In order to allow for arbitrary parameters the functions are defined by the following data types:

```
type gsl_multifit_nlinear_fdf
```

This data type defines a general system of functions with arbitrary parameters, the corresponding Jacobian matrix of derivatives, and optionally the second directional derivative of the functions for geodesic acceleration.

```
int (* f) (const gsl_vector * x, void * params, gsl_vector * f)
```

This function should store the n components of the vector $f(x)$ in `f` for argument `x` and arbitrary parameters `params`, returning an appropriate error code if the function cannot be computed.

```
int (* df) (const gsl_vector * x, void * params, gsl_matrix * J)
```

This function should store the n -by- p matrix result

$$J_{ij} = \partial f_i(x) / \partial x_j$$

in `J` for argument `x` and arbitrary parameters `params`, returning an appropriate error code if the matrix cannot be computed. If an analytic Jacobian is unavailable, or too expensive to compute, this function pointer may be set to `NULL`, in which case the Jacobian will be internally computed using finite difference approximations of the function `f`.

```
int (* fvv) (const gsl_vector * x, const gsl_vector * v, void * params, gsl_vector * fvv)
```

When geodesic acceleration is enabled, this function should store the n components of the vector $f_{vv}(x) = \sum_{\alpha\beta} v_\alpha v_\beta \frac{\partial}{\partial x_\alpha} \frac{\partial}{\partial x_\beta} f(x)$, representing second directional derivatives of the function to be minimized, into the output `fvv`. The parameter vector is provided in `x` and the velocity vector is provided in `v`, both of which have p components. The arbitrary parameters are given in `params`. If analytic expressions for $f_{vv}(x)$ are unavailable or too difficult to compute, this function pointer may be set to `NULL`, in which case $f_{vv}(x)$ will be computed internally using a finite difference approximation.

```
size_t n
```

the number of functions, i.e. the number of components of the vector `f`.

```
size_t p
```

the number of independent variables, i.e. the number of components of the vector `x`.

```
void * params
```

a pointer to the arbitrary parameters of the function.

```
size_t nevalf
```

This does not need to be set by the user. It counts the number of function evaluations and is initialized by the `_init` function.

```
size_t nevaldf
```

This does not need to be set by the user. It counts the number of Jacobian evaluations and is initialized by the `_init` function.

```
size_t nevalfvv
```

This does not need to be set by the user. It counts the number of $f_{vv}(x)$ evaluations and is initialized by the `_init` function.

type gsl_multilarge_nlinear_fdf

This data type defines a general system of functions with arbitrary parameters, a function to compute Ju or $J^T u$ for a given vector u , the normal equations matrix $J^T J$, and optionally the second directional derivative of the functions for geodesic acceleration.

```
int (* f) (const gsl_vector * x, void * params, gsl_vector * f)
```

This function should store the n components of the vector $f(x)$ in `f` for argument `x` and arbitrary parameters `params`, returning an appropriate error code if the function cannot be computed.

```
int (* df) (CBLAS_TRANSPOSE_t TransJ, const gsl_vector * x, const gsl_vector * u, void * params, gsl_vector * v, gs
```

If `TransJ` is equal to `CblasNoTrans`, then this function should compute the matrix-vector product Ju and store the result in `v`. If `TransJ` is equal to `CblasTrans`, then this function should compute the matrix-vector product $J^T u$ and store the result in `v`. Additionally, the normal equations matrix $J^T J$ should be stored in the lower half of `JTJ`. The input matrix `JTJ` could be set to `NULL`, for example by iterative methods which do not require this matrix, so the user should check for this prior to constructing the matrix. The input `params` contains the arbitrary parameters.

```
int (* fvv) (const gsl_vector * x, const gsl_vector * v, void * params, gsl_vector * fvv)
```

When geodesic acceleration is enabled, this function should store the n components of the vector $f_{vv}(x) = \sum_{\alpha\beta} v_\alpha v_\beta \frac{\partial}{\partial x_\alpha} \frac{\partial}{\partial x_\beta} f(x)$, representing second directional derivatives of the function to be minimized, into the output `fvv`. The parameter vector is provided in `x` and the velocity vector is provided in `v`, both of which have p components. The arbitrary parameters are given in `params`. If analytic expressions for $f_{vv}(x)$ are unavailable or too difficult to compute, this function pointer may be set to `NULL`, in which case $f_{vv}(x)$ will be computed internally using a finite difference approximation.

```
size_t n
```

the number of functions, i.e. the number of components of the vector `f`.

```
size_t p
```

the number of independent variables, i.e. the number of components of the vector `x`.

```
void * params
```

a pointer to the arbitrary parameters of the function.

```
size_t nevalf
```

This does not need to be set by the user. It counts the number of function evaluations and is initialized by the `_init` function.

```
size_t nevaldfu
```

This does not need to be set by the user. It counts the number of Jacobian matrix-vector evaluations (Ju or $J^T u$) and is initialized by the `_init` function.

```
size_t nevaldf2
```

This does not need to be set by the user. It counts the number of $J^T J$ evaluations and is initialized by the `_init` function.

```
size_t nevalfvv
```

This does not need to be set by the user. It counts the number of $f_{vv}(x)$ evaluations and is initialized by the `_init` function.

Note that when fitting a non-linear model against experimental data, the data is passed to the functions above using the `params` argument and the trial best-fit parameters through the `x` argument.

Iteration

The following functions drive the iteration of each **algorithm**. Each function performs one iteration of the trust region method and updates the state of the solver.

```
int gsl_multifit_nlinear_iterate(gsl_multifit_nlinear_workspace *w)
```

```
int gsl_multilarge_nlinear_iterate(gsl_multilarge_nlinear_workspace *w)
```

These functions perform a single iteration of the solver `w`. If the iteration encounters an unexpected problem then an error code will be returned. The solver workspace maintains a current estimate of the best-fit parameters at all times.

The solver workspace `w` contains the following entries, which can be used to track the progress of the solution:

```
gsl_vector * x
```

The current position, length p .

```
gsl_vector * f
```

The function residual vector at the current position $f(x)$, length n .

```
gsl_matrix * J
```

The Jacobian matrix at the current position $J(x)$, size n -by- p (only for `gsl_multifit_nlinear` interface).

```
gsl_vector * dx
```

The difference between the current position and the previous position, i.e. the last step δ , taken as a vector, length p .

These quantities can be accessed with the following functions,

```
gsl_vector *gsl_multifit_nlinear_position(const gsl_multifit_nlinear_workspace *w)
```

```
gsl_vector *gsl_multilarge_nlinear_position(const gsl_multilarge_nlinear_workspace *w)
```

These functions return the current position x (i.e. best-fit parameters) of the solver `w`.

```
gsl_vector *gsl_multifit_nlinear_residual(const gsl_multifit_nlinear_workspace *w)
```

```
gsl_vector *gsl_multilarge_nlinear_residual(const gsl_multilarge_nlinear_workspace *w)
```

These functions return the current residual vector $f(x)$ of the solver `w`. For weighted systems, the residual vector includes the weighting factor \sqrt{W} .

```
gsl_matrix *gsl_multifit_nlinear_jac(const gsl_multifit_nlinear_workspace *w)
```

This function returns a pointer to the n -by- p Jacobian matrix for the current iteration of the solver `w`. This function is available only for the `gsl_multifit_nlinear` interface.

```
size_t gsl_multifit_nlinear_niter(const gsl_multifit_nlinear_workspace *w)
```

```
size_t gsl_multilarge_nlinear_niter(const gsl_multilarge_nlinear_workspace *w)
```

These functions return the number of iterations performed so far. The iteration counter is updated on each call to the `_iterate` functions above, and reset to 0 in the `_init` functions.

```
int gsl_multifit_nlinear_rcond(double *rcond, const gsl_multifit_nlinear_workspace *w)
```

```
int gsl_multilarge_nlinear_rcond(double *rcond, const gsl_multilarge_nlinear_workspace *w)
```

This function estimates the reciprocal condition number of the Jacobian matrix at the current position x and stores it in `rcond`. The computed value is only an estimate to give the user a guideline as to the conditioning of their particular problem. Its calculation is based on which factorization method is used (Cholesky, QR, or SVD).

- For the Cholesky solver, the matrix $J^T J$ is factored at each iteration. Therefore this function will estimate the 1-norm condition number $rcond^2 = 1/(\|J^T J\|_1 \cdot \|(J^T J)^{-1}\|_1)$
- For the QR solver, J is factored as $J = QR$ at each iteration. For simplicity, this function calculates the 1-norm conditioning of only the R factor, $rcond = 1/(\|R\|_1 \cdot \|R^{-1}\|_1)$. This can be computed efficiently since R is upper triangular.
- For the SVD solver, in order to efficiently solve the trust region subproblem, the matrix which is factored is $J D^{-1}$, instead of J itself. The resulting singular values are used to provide the 2-norm reciprocal condition number, as $rcond = \sigma_{min} / \sigma_{max}$. Note that when using Moré scaling, $D \neq I$ and the resulting `rcond` estimate may be significantly different from the true `rcond` of J itself.

```
double gsl_multifit_nlinear_avratio(const gsl_multifit_nlinear_workspace *w)
```

```
double gsl_multilarge_nlinear_avratio(const gsl_multilarge_nlinear_workspace *w)
```

This function returns the current ratio $|a|/|v|$ of the acceleration correction term to the velocity step term. The acceleration term is computed only by the `gsl_multifit_nlinear_trs_lmaccel` and `gsl_multilarge_nlinear_trs_lmaccel` methods, so this ratio will be zero for other TRS methods.

Testing for Convergence

A minimization procedure should stop when one of the following conditions is true:

- A minimum has been found to within the user-specified precision.
- A user-specified maximum number of iterations has been reached.
- An error has occurred.

The handling of these conditions is under user control. The functions below allow the user to test the current estimate of the best-fit parameters in several standard ways.

```
int gsl_multifit_nlinear_test(const double xtol, const double gtol, const double ftol, int *info, const gsl_multifit_nlinear_workspace *w)
```

```
int gsl_multilarge_nlinear_test(const double xtol, const double gtol, const double ftol, int *info, const gsl_multilarge_nlinear_workspace *w)
```

These functions test for convergence of the minimization method using the following criteria:

- Testing for a small step size relative to the current parameter vector

$$|\delta_i| \leq xtol(|x_i| + xtol)$$

for each $0 \leq i < p$. Each element of the step vector δ is tested individually in case the different parameters have widely different scales. Adding `xtol` to $|x_i|$ helps the test avoid breaking down in situations where the true solution value $x_i = 0$. If this test succeeds, `info` is set to 1 and the function returns `GSL_SUCCESS`.

A general guideline for selecting the step tolerance is to choose $xtol = 10^{-d}$ where d is the number of accurate decimal digits desired in the solution x . See Dennis and Schnabel for more information.

- Testing for a small gradient ($g = \nabla\Phi(x) = J^T f$) indicating a local function minimum:

$$\max_i |g_i \times \max(x_i, 1)| \leq gtol \times \max(\Phi(x), 1)$$

This expression tests whether the ratio $(\nabla\Phi)_i x_i / \Phi$ is small. Testing this scaled gradient is a better than $\nabla\Phi$ alone since it is a dimensionless quantity and so independent of the scale of the problem. The `max` arguments help ensure the test doesn't break down in regions where x_i or $\Phi(x)$ are close to 0. If this test succeeds, `info` is set to 2 and the function returns `GSL_SUCCESS`.

A general guideline for choosing the gradient tolerance is to set

`gtol = GSL_DBL_EPSILON^(1/3)`. See Dennis and Schnabel for more information.

If none of the tests succeed, `info` is set to 0 and the function returns `GSL_CONTINUE`, indicating further iterations are required.

High Level Driver

These routines provide a high level wrapper that combines the iteration and convergence testing for easy use.

```
int gsl_multifit_nlinear_driver(const size_t maxiter, const double xtol, const double gtol, const double ftol,
void (*callback)(const size_t iter, void *params, const gsl_multifit_linear_workspace *w), void
*callback_params, int *info, gsl_multifit_nlinear_workspace *w)
```

```
int gsl_multilarge_nlinear_driver(const size_t maxiter, const double xtol, const double gtol, const double
ftol, void (*callback)(const size_t iter, void *params, const gsl_multilarge_linear_workspace *w), void
*callback_params, int *info, gsl_multilarge_nlinear_workspace *w)
```

These functions iterate the nonlinear least squares solver `w` for a maximum of `maxiter` iterations. After each iteration, the system is tested for convergence with the error tolerances `xtol`, `gtol` and `ftol`. Additionally, the user may supply a callback function `callback` which is called after each iteration, so that the user may save or print relevant quantities for each iteration. The parameter `callback_params` is passed to the `callback` function. The parameters `callback` and `callback_params` may be set to `NULL` to disable this feature. Upon successful convergence, the function returns `GSL_SUCCESS` and sets `info` to the reason for convergence (see `gsl_multifit_nlinear_test()`). If the function has not converged after `maxiter` iterations, `GSL_EMAXITER` is returned. In rare cases, during an iteration the `algorithm` may be unable to find a new acceptable step δ to take. In this case, `GSL_ENOPROG` is returned indicating no further progress can be made. If your problem is having difficulty converging, see [Troubleshooting](#) for further guidance.

Covariance matrix of best fit parameters

```
int gsl_multifit_nlinear_covar(const gsl_matrix *J, const double epsrel, gsl_matrix *covar)
```

```
int gsl_multilarge_nlinear_covar(gsl_matrix *covar, gsl_multilarge_nlinear_workspace *w)
```

This function computes the covariance matrix of best-fit parameters using the Jacobian matrix `J` and stores it in `covar`. The parameter `epsrel` is used to remove linear-dependent columns when `J` is rank deficient.

The covariance matrix is given by,

$$C = (J^T J)^{-1}$$

or in the weighted case,

$$C = (J^T W J)^{-1}$$

and is computed using the factored form of the Jacobian (Cholesky, QR, or SVD). Any columns of R which satisfy

$$|R_{kk}| \leq \epsilon_{\text{psrel}} |R_{11}|$$

are considered linearly-dependent and are excluded from the covariance matrix (the corresponding rows and columns of the covariance matrix are set to zero).

If the minimisation uses the weighted least-squares function $f_i = (Y(x, t_i) - y_i)/\sigma_i$ then the covariance matrix above gives the statistical error on the best-fit parameters resulting from the Gaussian errors σ_i on the underlying data y_i . This can be verified from the relation $\delta f = J \delta c$ and the fact that the fluctuations in f from the data y_i are normalised by σ_i and so satisfy

$$\langle \delta f \delta f^T \rangle = I$$

For an unweighted least-squares function $f_i = (Y(x, t_i) - y_i)$ the covariance matrix above should be multiplied by the variance of the residuals about the best-fit $\sigma^2 = \sum (y_i - Y(x, t_i))^2 / (n - p)$ to give the variance-covariance matrix $\sigma^2 C$. This estimates the statistical error on the best-fit parameters from the scatter of the underlying data.

For more information about covariance matrices see [Linear Least-Squares Overview](#).

Troubleshooting

When developing a code to solve a nonlinear least squares problem, here are a few considerations to keep in mind.

1. The most common difficulty is the accurate implementation of the Jacobian matrix. If the analytic Jacobian is not properly provided to the solver, this can hinder and many times prevent convergence of the method. When developing a new nonlinear least squares code, it often helps to compare the program output with the internally computed finite difference Jacobian and the user supplied analytic Jacobian. If there is a large difference in coefficients, it is likely the analytic Jacobian is incorrectly implemented.
2. If your code is having difficulty converging, the next thing to check is the starting point provided to the solver. The methods of this chapter are local methods, meaning if you provide a starting point far away from the true minimum, the method may converge to a local minimum or not converge at all. Sometimes it is possible to solve a linearized approximation to the nonlinear problem, and use the linear solution as the starting point to the nonlinear problem.
3. If the various parameters of the coefficient vector x vary widely in magnitude, then the problem is said to be badly scaled. The methods of this chapter do attempt to automatically rescale the elements of x to have roughly the same order of magnitude, but in extreme cases this could still cause problems for convergence. In these cases it is recommended for the user to scale their parameter vector x so that each parameter spans roughly the same range, say $[-1, 1]$. The solution vector can be backscaled to recover the original units of the problem.

Examples

The following example programs demonstrate the nonlinear least squares fitting capabilities.

Exponential Fitting Example

The following example program fits a weighted exponential model with background to experimental data, $Y = A \exp(-\lambda t) + b$. The first part of the program sets up the functions `expb_f()` and `expb_df()` to calculate the model and its Jacobian. The appropriate fitting function is given by,

$$f_i = (A \exp(-\lambda t_i) + b) - y_i$$

where we have chosen $t_i = iT/(N - 1)$, where N is the number of data points fitted, so that $t_i \in [0, T]$. The Jacobian matrix J is the derivative of these functions with respect to the three parameters (A, λ, b) . It is given by,

$$J_{ij} = \frac{\partial f_i}{\partial x_j}$$

where $x_0 = A$, $x_1 = \lambda$ and $x_2 = b$. The i -th row of the Jacobian is therefore

$$J_i = \begin{pmatrix} \exp(-\lambda t_i) & -t_i A \exp(-\lambda t_i) & 1 \end{pmatrix}$$

The main part of the program sets up a **Levenberg-Marquardt** solver and some simulated random data. The data uses the known parameters (5.0,1.5,1.0) combined with Gaussian noise (standard deviation = 0.1) with a maximum time $T = 3$ and $N = 100$ timesteps. The initial guess for the parameters is chosen as (1.0, 1.0, 0.0). The iteration terminates when the relative change in x is smaller than 10^{-8} , or when the magnitude of the gradient falls below 10^{-8} . Here are the results of running the program:

```
iter 0: A = 1.0000, lambda = 1.0000, b = 0.0000, cond(J) = inf, |f(x)| = 88.4448
iter 1: A = 4.5109, lambda = 2.5258, b = 1.0704, cond(J) = 26.2686, |f(x)| = 24.0646
iter 2: A = 4.8565, lambda = 1.7442, b = 1.1669, cond(J) = 23.7470, |f(x)| = 11.9797
iter 3: A = 4.9356, lambda = 1.5713, b = 1.0767, cond(J) = 17.5849, |f(x)| = 10.7355
iter 4: A = 4.8678, lambda = 1.4838, b = 1.0252, cond(J) = 16.3428, |f(x)| = 10.5000
iter 5: A = 4.8118, lambda = 1.4481, b = 1.0076, cond(J) = 15.7925, |f(x)| = 10.4786
iter 6: A = 4.7983, lambda = 1.4404, b = 1.0041, cond(J) = 15.5840, |f(x)| = 10.4778
iter 7: A = 4.7967, lambda = 1.4395, b = 1.0037, cond(J) = 15.5396, |f(x)| = 10.4778
iter 8: A = 4.7965, lambda = 1.4394, b = 1.0037, cond(J) = 15.5344, |f(x)| = 10.4778
iter 9: A = 4.7965, lambda = 1.4394, b = 1.0037, cond(J) = 15.5339, |f(x)| = 10.4778
iter 10: A = 4.7965, lambda = 1.4394, b = 1.0037, cond(J) = 15.5339, |f(x)| = 10.4778
iter 11: A = 4.7965, lambda = 1.4394, b = 1.0037, cond(J) = 15.5339, |f(x)| = 10.4778
summary from method 'trust-region/levenberg-marquardt'
number of iterations: 11
function evaluations: 16
Jacobian evaluations: 12
reason for stopping: small gradient
initial |f(x)| = 88.444756
final |f(x)| = 10.477801
chisq/dof = 1.1318
A = 4.79653 +/- 0.18704
lambda = 1.43937 +/- 0.07390
b = 1.00368 +/- 0.03473
status = success
```

The approximate values of the parameters are found correctly, and the chi-squared value indicates a good fit (the chi-squared per degree of freedom is approximately 1). In this case the errors on the parameters can be estimated from the square roots of the diagonal elements of the covariance matrix. If the chi-squared value shows a poor fit (i.e. $\chi^2/(n - p) \gg 1$) then the error estimates obtained from the covariance matrix will be too small. In the example program the error estimates are multiplied by $\sqrt{\chi^2/(n - p)}$ in this case, a common way of increasing the errors for a poor fit. Note that a poor fit will result from the use of an inappropriate model, and the scaled error estimates may then be outside the range of validity for Gaussian errors.

Additionally, we see that the condition number of $J(x)$ stays reasonably small throughout the iteration. This indicates we could safely switch to the Cholesky solver for speed improvement, although this particular system is too small to really benefit.

Fig. 35 shows the fitted curve with the original data.




Fig. 35 Exponential fitted curve with data

```

#include <stdlib.h>
#include <stdio.h>
#include <gsl/gsl_rng.h>
#include <gsl/gsl_randist.h>
#include <gsl/gsl_matrix.h>
#include <gsl/gsl_vector.h>
#include <gsl/gsl_blas.h>
#include <gsl/gsl_multifit_nlinear.h>

#define N      100    /* number of data points to fit */
#define TMAX   (3.0)  /* time variable in [0,TMAX] */

struct data {
    size_t n;
    double * t;
    double * y;
};

int
expb_f (const gsl_vector * x, void *data,
        gsl_vector * f)
{
    struct data * data = (struct data *)data;

    size_t n = ((struct data *)data)->n;
    double *t = ((struct data *)data)->t;
    double *y = ((struct data *)data)->y;

    double A = gsl_vector_get (x, 0);
    double lambda = gsl_vector_get (x, 1);
    double b = gsl_vector_get (x, 2);

    size_t i;

    for (i = 0; i < n; i++)
    {
        /* Model  $Y_i = A * \exp(-\lambda * t_i) + b$  */
        double Yi = A * exp (-lambda * t[i]) + b;
        gsl_vector_set (f, i, Yi - y[i]);
    }

    return GSL_SUCCESS;
}

int
expb_df (const gsl_vector * x, void *data,
         gsl_matrix * J)
{
    struct data * data = (struct data *)data;

    size_t n = ((struct data *)data)->n;
    double *t = ((struct data *)data)->t;

    double A = gsl_vector_get (x, 0);
    double lambda = gsl_vector_get (x, 1);

    size_t i;

    for (i = 0; i < n; i++)
    {
        /* Jacobian matrix  $J(i,j) = df_i / dx_j$ , */
        /* where  $fi = (Y_i - y_i)/\sigma[i]$ , */
        /*  $Y_i = A * \exp(-\lambda * t_i) + b$  */
        /* and the  $x_j$  are the parameters (A, lambda, b) */
        double e = exp(-lambda * t[i]);
        gsl_matrix_set (J, i, 0, e);
        gsl_matrix_set (J, i, 1, -t[i] * A * e);
        gsl_matrix_set (J, i, 2, 1.0);
    }

    return GSL_SUCCESS;
}

void
callback(const size_t iter, void *params,
         const gsl_multifit_nlinear_workspace *w)
{
    gsl_vector *f = gsl_multifit_nlinear_residual(w);
    gsl_vector *x = gsl_multifit_nlinear_position(w);
    double rcond;

    /* compute reciprocal condition number of J(x) */

```

```

gsl_multifit_nlinear_rcond(&rcond, w);

    fprintf(stderr, "iter %2zu: A = %.4f, lambda = %.4f, b = %.4f, cond(J) = %8.4f, |f(x)| = %.4f\n",
        iter,
        gsl_vector_get(x, 0),
        gsl_vector_get(x, 1),
        gsl_vector_get(x, 2),
        1.0 / rcond,
        gsl_blas_dnorm2(f));
}

int
main (void)
{
    const gsl_multifit_nlinear_type *T = gsl_multifit_nlinear_trust;
    gsl_multifit_nlinear_workspace *w;
    gsl_multifit_nlinear_fdf fdf;
    gsl_multifit_nlinear_parameters fdf_params =
        gsl_multifit_nlinear_default_parameters();
    const size_t n = N;
    const size_t p = 3;

    gsl_vector *f;
    gsl_matrix *J;
    gsl_matrix *covar = gsl_matrix_alloc (p, p);
    double t[N], y[N], weights[N];
    struct data d = { n, t, y };
    double x_init[3] = { 1.0, 1.0, 0.0 }; /* starting values */
    gsl_vector_view x = gsl_vector_view_array (x_init, p);
    gsl_vector_view wts = gsl_vector_view_array(weights, n);
    gsl_rng * r;
    double chisq, chisq0;
    int status, info;
    size_t i;

    const double xtol = 1e-8;
    const double gtol = 1e-8;
    const double ftol = 0.0;

    gsl_rng_env_setup();
    r = gsl_rng_alloc(gsl_rng_default);

    /* define the function to be minimized */
    fdf.f = expb_f;
    fdf.df = expb_df; /* set to NULL for finite-difference Jacobian */
    fdf.fvv = NULL; /* not using geodesic acceleration */
    fdf.n = n;
    fdf.p = p;
    fdf.params = &d;

    /* this is the data to be fitted */
    for (i = 0; i < n; i++)
    {
        double ti = i * TMAX / (n - 1.0);
        double yi = 1.0 + 5 * exp (-1.5 * ti);
        double si = 0.1 * yi;
        double dy = gsl_ran_gaussian(r, si);

        t[i] = ti;
        y[i] = yi + dy;
        weights[i] = 1.0 / (si * si);
        printf ("data: %g %g %g\n", ti, y[i], si);
    };

    /* allocate workspace with default parameters */
    w = gsl_multifit_nlinear_alloc (T, &fdf_params, n, p);

    /* initialize solver with starting point and weights */
    gsl_multifit_nlinear_winit (&x.vector, &wts.vector, &fdf, w);

    /* compute initial cost function */
    f = gsl_multifit_nlinear_residual(w);
    gsl_blas_ddot(f, f, &chisq0);

    /* solve the system with a maximum of 100 iterations */
    status = gsl_multifit_nlinear_driver(100, xtol, gtol, ftol,
        callback, NULL, &info, w);
}

```

```

/* compute covariance of best fit parameters */
J = gsl_multifit_nlinear_jac(w);
gsl_multifit_nlinear_covar (J, 0.0, covar);

/* compute final cost */
gsl_blas_ddot(f, f, &chisq);

#define FIT(i) gsl_vector_get(w->x, i)
#define ERR(i) sqrt(gsl_matrix_get(covar,i,i))

    fprintf(stderr, "summary from method '%s/%s'\n",
        gsl_multifit_nlinear_name(w),
        gsl_multifit_nlinear_trs_name(w));
    fprintf(stderr, "number of iterations: %zu\n",
        gsl_multifit_nlinear_niter(w));
    fprintf(stderr, "function evaluations: %zu\n", fdf.nevalf);
    fprintf(stderr, "Jacobian evaluations: %zu\n", fdf.nevaldf);
    fprintf(stderr, "reason for stopping: %s\n",
        (info == 1) ? "small step size" : "small gradient");
    fprintf(stderr, "initial |f(x)| = %f\n", sqrt(chisq0));
    fprintf(stderr, "final   |f(x)| = %f\n", sqrt(chisq));

    {
        double dof = n - p;
        double c = GSL_MAX_DBL(1, sqrt(chisq / dof));

        fprintf(stderr, "chisq/dof = %g\n", chisq / dof);

        fprintf (stderr, "A      = %.5f +/- %.5f\n", FIT(0), c*ERR(0));
        fprintf (stderr, "lambda = %.5f +/- %.5f\n", FIT(1), c*ERR(1));
        fprintf (stderr, "b      = %.5f +/- %.5f\n", FIT(2), c*ERR(2));
    }

    fprintf (stderr, "status = %s\n", gsl_strerror (status));

    gsl_multifit_nlinear_free (w);
    gsl_matrix_free (covar);
    gsl_rng_free (r);

    return 0;
}

```

Geodesic Acceleration Example 1

The following example program minimizes a modified Rosenbrock function, which is characterized by a narrow canyon with steep walls. The starting point is selected high on the canyon wall, so the solver must first find the canyon bottom and then navigate to the minimum. The problem is solved both with and without using geodesic acceleration for comparison. The cost function is given by

$$\begin{aligned}\Phi(x) &= \frac{1}{2}(f_1^2 + f_2^2) \\ f_1 &= 100(x_2 - x_1^2) \\ f_2 &= 1 - x_1\end{aligned}$$

The Jacobian matrix is

$$J = \begin{pmatrix} \frac{\partial f_1}{\partial x_1} & \frac{\partial f_1}{\partial x_2} \\ \frac{\partial f_2}{\partial x_1} & \frac{\partial f_2}{\partial x_2} \end{pmatrix} = \begin{pmatrix} -200x_1 & 100 \\ -1 & 0 \end{pmatrix}$$

In order to use geodesic acceleration, the user must provide the second directional derivative of each residual in the velocity direction, $D_v^2 f_i = \sum_{\alpha\beta} v_\alpha v_\beta \partial_\alpha \partial_\beta f_i$. The velocity vector v is provided by the solver. For this example, these derivatives are

$$f_{vv} = D_v^2 \begin{pmatrix} f_1 \\ f_2 \end{pmatrix} = \begin{pmatrix} -200v_1^2 \\ 0 \end{pmatrix}$$

The solution of this minimization problem is

$$x^* = \begin{pmatrix} 1 \\ 1 \end{pmatrix}$$

$$\Phi(x^*) = 0$$

The program output is shown below:

```

=== Solving system without acceleration ===
NITER      = 53
NFEV       = 56
NJEV       = 54
NAEV       = 0
initial cost = 2.250225000000e+04
final cost   = 6.674986031430e-18
final x      = (9.999999974165e-01, 9.999999948328e-01)
final cond(J) = 6.000096055094e+02

=== Solving system with acceleration ===
NITER      = 15
NFEV       = 17
NJEV       = 16
NAEV       = 16
initial cost = 2.250225000000e+04
final cost   = 7.518932873279e-19
final x      = (9.999999991329e-01, 9.999999982657e-01)
final cond(J) = 6.000097233278e+02

```




Fig. 36 Paths taken by solver for Rosenbrock function

We can see that enabling geodesic acceleration requires less than a third of the number of Jacobian evaluations in order to locate the minimum. The path taken by both methods is shown in Fig. 36. The contours show the cost function $\Phi(x_1, x_2)$. We see that both methods quickly find the canyon bottom, but the geodesic acceleration method navigates along the bottom to the solution with significantly fewer iterations.

The program is given below.

```

#include <stdlib.h>
#include <stdio.h>
#include <gsl/gsl_vector.h>
#include <gsl/gsl_matrix.h>
#include <gsl/gsl_blas.h>
#include <gsl/gsl_multifit_nlinear.h>

int
func_f (const gsl_vector * x, void *params, gsl_vector * f)
{
    double x1 = gsl_vector_get(x, 0);
    double x2 = gsl_vector_get(x, 1);

    gsl_vector_set(f, 0, 100.0 * (x2 - x1*x1));
    gsl_vector_set(f, 1, 1.0 - x1);

    return GSL_SUCCESS;
}

int
func_df (const gsl_vector * x, void *params, gsl_matrix * J)
{
    double x1 = gsl_vector_get(x, 0);

    gsl_matrix_set(J, 0, 0, -200.0*x1);
    gsl_matrix_set(J, 0, 1, 100.0);
    gsl_matrix_set(J, 1, 0, -1.0);
    gsl_matrix_set(J, 1, 1, 0.0);

    return GSL_SUCCESS;
}

int
func_fvv (const gsl_vector * x, const gsl_vector * v,
          void *params, gsl_vector * fvv)
{
    double v1 = gsl_vector_get(v, 0);

    gsl_vector_set(fvv, 0, -200.0 * v1 * v1);
    gsl_vector_set(fvv, 1, 0.0);

    return GSL_SUCCESS;
}

void
callback(const size_t iter, void *params,
          const gsl_multifit_nlinear_workspace *w)
{
    gsl_vector * x = gsl_multifit_nlinear_position(w);

    /* print out current location */
    printf("%f %f\n",
           gsl_vector_get(x, 0),
           gsl_vector_get(x, 1));
}

void
solve_system(gsl_vector *x0, gsl_multifit_nlinear_fdf *fdf,
             gsl_multifit_nlinear_parameters *params)
{
    const gsl_multifit_nlinear_type *T = gsl_multifit_nlinear_trust;
    const size_t max_iter = 200;
    const double xtol = 1.0e-8;
    const double gtol = 1.0e-8;
    const double ftol = 1.0e-8;
    const size_t n = fdf->n;
    const size_t p = fdf->p;
    gsl_multifit_nlinear_workspace *work =
        gsl_multifit_nlinear_alloc(T, params, n, p);
    gsl_vector * f = gsl_multifit_nlinear_residual(work);
    gsl_vector * x = gsl_multifit_nlinear_position(work);
    int info;
    double chisq0, chisq, rcond;

    /* initialize solver */
    gsl_multifit_nlinear_init(x0, fdf, work);

    /* store initial cost */

```

```

gsl_blas_ddot(f, f, &chisq0);

/* iterate until convergence */
gsl_multifit_nlinear_driver(max_iter, xtol, gtol, ftol,
                           callback, NULL, &info, work);

/* store final cost */
gsl_blas_ddot(f, f, &chisq);

/* store cond(J(x)) */
gsl_multifit_nlinear_rcond(&rcond, work);

/* print summary */

fprintf(stderr, "NITER      = %zu\n", gsl_multifit_nlinear_niter(work));
fprintf(stderr, "NFEV      = %zu\n", fdf->nevalf);
fprintf(stderr, "NJEV      = %zu\n", fdf->nevaldf);
fprintf(stderr, "NAEV      = %zu\n", fdf->nevalfvv);
fprintf(stderr, "initial cost = %.12e\n", chisq0);
fprintf(stderr, "final cost  = %.12e\n", chisq);
fprintf(stderr, "final x     = (%.12e, %.12e)\n",
        gsl_vector_get(x, 0), gsl_vector_get(x, 1));
fprintf(stderr, "final cond(J) = %.12e\n", 1.0 / rcond);

printf("\n\n");

gsl_multifit_nlinear_free(work);
}

int
main (void)
{
    const size_t n = 2;
    const size_t p = 2;
    gsl_vector *f = gsl_vector_alloc(n);
    gsl_vector *x = gsl_vector_alloc(p);
    gsl_multifit_nlinear_fdf fdf;
    gsl_multifit_nlinear_parameters fdf_params =
        gsl_multifit_nlinear_default_parameters();

    /* print map of Phi(x1, x2) */
    {
        double x1, x2, chisq;
        double *f1 = gsl_vector_ptr(f, 0);
        double *f2 = gsl_vector_ptr(f, 1);

        for (x1 = -1.2; x1 < 1.3; x1 += 0.1)
        {
            for (x2 = -0.5; x2 < 2.1; x2 += 0.1)
            {
                gsl_vector_set(x, 0, x1);
                gsl_vector_set(x, 1, x2);
                func_f(x, NULL, f);

                chisq = (*f1) * (*f1) + (*f2) * (*f2);
                printf("%f %f %f\n", x1, x2, chisq);
            }
            printf("\n");
        }
        printf("\n\n");
    }

    /* define function to be minimized */
    fdf.f = func_f;
    fdf.df = func_df;
    fdf.fvv = func_fvv;
    fdf.n = n;
    fdf.p = p;
    fdf.params = NULL;

    /* starting point */
    gsl_vector_set(x, 0, -0.5);
    gsl_vector_set(x, 1, 1.75);

    fprintf(stderr, "=== Solving system without acceleration ===\n");
    fdf_params.trs = gsl_multifit_nlinear_trs_lm;
    solve_system(x, &fdf, &fdf_params);

    fprintf(stderr, "=== Solving system with acceleration ===\n");

```



```

fdf_params.trs = gsl_multifit_nlinear_trs_lmaccel;
solve_system(x, &fdf, &fdf_params);

gsl_vector_free(f);
gsl_vector_free(x);

return 0;
}

```

Geodesic Acceleration Example 2

The following example fits a set of data to a Gaussian model using the **Levenberg-Marquardt** method with geodesic acceleration. The cost function is

$$\Phi(x) = \frac{1}{2} \sum_i f_i^2$$

$$f_i = y_i - Y(a, b, c; t_i)$$

where y_i is the measured data point at time t_i , and the model is specified by

$$Y(a, b, c; t) = a \exp \left[-\frac{1}{2} \left(\frac{t - b}{c} \right)^2 \right]$$

The parameters a, b, c represent the amplitude, mean, and width of the Gaussian respectively. The program below generates the y_i data on $[0, 1]$ using the values $a = 5, b = 0.4, c = 0.15$ and adding random noise. The i -th row of the Jacobian is

$$J_{i,:} = \left(\frac{\partial f_i}{\partial a} \quad \frac{\partial f_i}{\partial b} \quad \frac{\partial f_i}{\partial c} \right) = \left(-e_i \quad -\frac{a}{c} z_i e_i \quad -\frac{a}{c} z_i^2 e_i \right)$$

where

$$z_i = \frac{t_i - b}{c}$$

$$e_i = \exp \left(-\frac{1}{2} z_i^2 \right)$$

In order to use geodesic acceleration, we need the second directional derivative of the residuals in the velocity direction, $D_v^2 f_i = \sum_{\alpha\beta} v_\alpha v_\beta \partial_\alpha \partial_\beta f_i$, where v is provided by the solver. To compute this, it is helpful to make a table of all second derivatives of the residuals f_i with respect to each combination of model parameters. This table is

	$\frac{\partial}{\partial a}$	$\frac{\partial}{\partial b}$	$\frac{\partial}{\partial c}$
$\frac{\partial}{\partial a}$	0	$-\frac{z_i}{c} e_i$	$-\frac{z_i^2}{c} e_i$
$\frac{\partial}{\partial b}$	$\frac{a}{c^2} (1 - z_i^2) e_i$	$\frac{a}{c^2} z_i (2 - z_i^2) e_i$	$\frac{a}{c^2} z_i^2 (3 - z_i^2) e_i$
$\frac{\partial}{\partial c}$			

The lower half of the table is omitted since it is symmetric. Then, the second directional derivative of f_i is

$$D_v^2 f_i = v_a^2 \partial_a^2 f_i + 2v_a v_b \partial_a \partial_b f_i + 2v_a v_c \partial_a \partial_c f_i + v_b^2 \partial_b^2 f_i + 2v_b v_c \partial_b \partial_c f_i + v_c^2 \partial_c^2 f_i$$

The factors of 2 come from the symmetry of the mixed second partial derivatives. The iteration is started using the initial guess $a = 1, b = 0, c = 1$. The program output is shown below:

```

iter 0: a = 1.0000, b = 0.0000, c = 1.0000, |a|/|v| = 0.0000 cond(J) =      inf, |f(x)| =
35.4785
iter 1: a = 1.5708, b = 0.5321, c = 0.5219, |a|/|v| = 0.3093 cond(J) = 29.0443, |f(x)| =
31.1042
iter 2: a = 1.7387, b = 0.4040, c = 0.4568, |a|/|v| = 0.1199 cond(J) =  3.5256, |f(x)| =
28.7217
iter 3: a = 2.2340, b = 0.3829, c = 0.3053, |a|/|v| = 0.3308 cond(J) =  4.5121, |f(x)| =
23.8074
iter 4: a = 3.2275, b = 0.3952, c = 0.2243, |a|/|v| = 0.2784 cond(J) =  8.6499, |f(x)| =
15.6003
iter 5: a = 4.3347, b = 0.3974, c = 0.1752, |a|/|v| = 0.2029 cond(J) = 15.1732, |f(x)| =
7.5908
iter 6: a = 4.9352, b = 0.3992, c = 0.1536, |a|/|v| = 0.1001 cond(J) = 26.6621, |f(x)| =
4.8402
iter 7: a = 5.0716, b = 0.3994, c = 0.1498, |a|/|v| = 0.0166 cond(J) = 34.6922, |f(x)| =
4.7103
iter 8: a = 5.0828, b = 0.3994, c = 0.1495, |a|/|v| = 0.0012 cond(J) = 36.5422, |f(x)| =
4.7095
iter 9: a = 5.0831, b = 0.3994, c = 0.1495, |a|/|v| = 0.0000 cond(J) = 36.6929, |f(x)| =
4.7095
iter 10: a = 5.0831, b = 0.3994, c = 0.1495, |a|/|v| = 0.0000 cond(J) = 36.6975, |f(x)| =
4.7095
iter 11: a = 5.0831, b = 0.3994, c = 0.1495, |a|/|v| = 0.0000 cond(J) = 36.6976, |f(x)| =
4.7095
NITER          = 11
NFEV           = 18
NJEV           = 12
NAEV           = 17
initial cost   = 1.258724737288e+03
final cost     = 2.217977560180e+01
final x        = (5.083101559156e+00, 3.994484109594e-01, 1.494898e-01)
final cond(J) = 3.669757713403e+01

```

We see the method converges after 11 iterations. For comparison the standard **Levenberg-Marquardt** method requires 26 iterations and so the Gaussian fitting problem benefits substantially from the geodesic acceleration correction. The column marked `|a|/|v|` above shows the ratio of the acceleration term to the velocity term as the iteration progresses. Larger values of this ratio indicate that the geodesic acceleration correction term is contributing substantial information to the solver relative to the standard LM velocity step.

The data and fitted model are shown in [Fig. 37](#).




Fig. 37 Gaussian model fitted to data

The program is given below.

```

#include <stdlib.h>
#include <stdio.h>
#include <gsl/gsl_vector.h>
#include <gsl/gsl_matrix.h>
#include <gsl/gsl_blas.h>
#include <gsl/gsl_multifit_nlinear.h>
#include <gsl/gsl_rng.h>
#include <gsl/gsl_randist.h>

struct data
{
    double *t;
    double *y;
    size_t n;
};

/* model function:  $a * \exp(-1/2 * [(t - b) / c]^2)$  */
double
gaussian(const double a, const double b, const double c, const double t)
{
    const double z = (t - b) / c;
    return (a * exp(-0.5 * z * z));
}

int
func_f (const gsl_vector * x, void *params, gsl_vector * f)
{
    struct data *d = (struct data *) params;
    double a = gsl_vector_get(x, 0);
    double b = gsl_vector_get(x, 1);
    double c = gsl_vector_get(x, 2);
    size_t i;

    for (i = 0; i < d->n; ++i)
    {
        double ti = d->t[i];
        double yi = d->y[i];
        double y = gaussian(a, b, c, ti);

        gsl_vector_set(f, i, yi - y);
    }

    return GSL_SUCCESS;
}

int
func_df (const gsl_vector * x, void *params, gsl_matrix * J)
{
    struct data *d = (struct data *) params;
    double a = gsl_vector_get(x, 0);
    double b = gsl_vector_get(x, 1);
    double c = gsl_vector_get(x, 2);
    size_t i;

    for (i = 0; i < d->n; ++i)
    {
        double ti = d->t[i];
        double zi = (ti - b) / c;
        double ei = exp(-0.5 * zi * zi);

        gsl_matrix_set(J, i, 0, -ei);
        gsl_matrix_set(J, i, 1, -(a / c) * ei * zi);
        gsl_matrix_set(J, i, 2, -(a / c) * ei * zi * zi);
    }

    return GSL_SUCCESS;
}

int
func_fvv (const gsl_vector * x, const gsl_vector * v,
          void *params, gsl_vector * fvv)
{
    struct data *d = (struct data *) params;
    double a = gsl_vector_get(x, 0);
    double b = gsl_vector_get(x, 1);
    double c = gsl_vector_get(x, 2);
    double va = gsl_vector_get(v, 0);
    double vb = gsl_vector_get(v, 1);

```

[illegible]

```

/* store final cost */
gsl_blas_ddot(f, f, &chisq);

/* store cond(J(x)) */
gsl_multifit_nlinear_rcond(&rcond, work);

gsl_vector_memcpy(x, y);

/* print summary */

fprintf(stderr, "NITER      = %zu\n", gsl_multifit_nlinear_niter(work));
fprintf(stderr, "NFEV      = %zu\n", fdf->nevalf);
fprintf(stderr, "NJEV      = %zu\n", fdf->nevaldf);
fprintf(stderr, "NAEV      = %zu\n", fdf->nevalfvv);
fprintf(stderr, "initial cost = %.12e\n", chisq0);
fprintf(stderr, "final cost  = %.12e\n", chisq);
fprintf(stderr, "final x    = (%.12e, %.12e, %.12e)\n",
        gsl_vector_get(x, 0), gsl_vector_get(x, 1), gsl_vector_get(x, 2));
fprintf(stderr, "final cond(J) = %.12e\n", 1.0 / rcond);

gsl_multifit_nlinear_free(work);
}

int
main (void)
{
    const size_t n = 300; /* number of data points to fit */
    const size_t p = 3;   /* number of model parameters */
    const double a = 5.0; /* amplitude */
    const double b = 0.4; /* center */
    const double c = 0.15; /* width */
    const gsl_rng_type * T = gsl_rng_default;
    gsl_vector *f = gsl_vector_alloc(n);
    gsl_vector *x = gsl_vector_alloc(p);
    gsl_multifit_nlinear_fdf fdf;
    gsl_multifit_nlinear_parameters fdf_params =
        gsl_multifit_nlinear_default_parameters();
    struct data fit_data;
    gsl_rng * r;
    size_t i;

    gsl_rng_env_setup ();
    r = gsl_rng_alloc (T);

    fit_data.t = malloc(n * sizeof(double));
    fit_data.y = malloc(n * sizeof(double));
    fit_data.n = n;

    /* generate synthetic data with noise */
    for (i = 0; i < n; ++i)
    {
        double t = (double)i / (double) n;
        double y0 = gaussian(a, b, c, t);
        double dy = gsl_ran_gaussian (r, 0.1 * y0);

        fit_data.t[i] = t;
        fit_data.y[i] = y0 + dy;
    }

    /* define function to be minimized */
    fdf.f = func_f;
    fdf.df = func_df;
    fdf.fvv = func_fvv;
    fdf.n = n;
    fdf.p = p;
    fdf.params = &fit_data;

    /* starting point */
    gsl_vector_set(x, 0, 1.0);
    gsl_vector_set(x, 1, 0.0);
    gsl_vector_set(x, 2, 1.0);

    fdf_params.trs = gsl_multifit_nlinear_trs_lmaccel;
    solve_system(x, &fdf, &fdf_params);

    /* print data and model */
    {
        double A = gsl_vector_get(x, 0);

```

```

double B = gsl_vector_get(x, 1);
double C = gsl_vector_get(x, 2);

for (i = 0; i < n; ++i)
{
    double ti = fit_data.t[i];
    double yi = fit_data.y[i];
    double fi = gaussian(A, B, C, ti);

    printf("%f %f %f\n", ti, yi, fi);
}

gsl_vector_free(f);
gsl_vector_free(x);
gsl_rng_free(r);

return 0;
}

```

Comparing TRS Methods Example

The following program compares all available nonlinear least squares trust-region subproblem (TRS) methods on the Branin function, a common optimization test problem. The cost function is

$$\Phi(x) = \frac{1}{2}(f_1^2 + f_2^2)$$

$$f_1 = x_2 + a_1 x_1^2 + a_2 x_1 + a_3$$

$$f_2 = \sqrt{a_4} \sqrt{1 + (1 - a_5) \cos x_1}$$

with $a_1 = -\frac{5.1}{4\pi^2}$, $a_2 = \frac{5}{\pi}$, $a_3 = -6$, $a_4 = 10$, $a_5 = \frac{1}{8\pi}$. There are three minima of this function in the range $(x_1, x_2) \in [-5, 15] \times [-5, 15]$. The program below uses the starting point $(x_1, x_2) = (6, 14.5)$ and calculates the solution with all available nonlinear least squares TRS methods. The program output is shown below:

Method	NITER	NFEV	NJEV	Initial Cost	Final cost	Final cond(J)	Final x
Levenberg-marquardt (-3.14e+00, 1.23e+01)	20	27	21	1.9874e+02	3.9789e-01	6.1399e+07	
Levenberg-marquardt+accel (3.14e+00, 2.27e+00)	27	36	28	1.9874e+02	3.9789e-01	1.4465e+07	
dogleg (3.14e+00, 2.28e+00)	23	64	23	1.9874e+02	3.9789e-01	5.0692e+08	
double-dogleg (3.14e+00, 2.27e+00)	24	69	24	1.9874e+02	3.9789e-01	3.4879e+07	
2D-subspace (3.14e+00, 2.27e+00)	23	54	24	1.9874e+02	3.9789e-01	2.5142e+07	

The first row of output above corresponds to standard **Levenberg-Marquardt**, while the second row includes geodesic acceleration. We see that the standard LM method converges to the minimum at $(-\pi, 12.275)$ and also uses the least number of iterations and Jacobian evaluations. All other methods converge to the minimum $(\pi, 2.275)$ and perform similarly in terms of number of Jacobian evaluations. We see that J is fairly ill-conditioned at both minima, indicating that the QR (or SVD) solver is the best choice for this problem. Since there are only two parameters in this optimization problem, we can easily visualize the paths taken by each method, which are shown in Fig. 38. The figure shows contours of the cost function $\Phi(x_1, x_2)$ which exhibits three global minima in the range $[-5, 15] \times [-5, 15]$. The paths taken by each solver are shown as colored lines.




Fig. 38 Paths taken for different TRS methods for the Branin function

The program is given below.


```

#include <stdlib.h>
#include <stdio.h>
#include <gsl/gsl_vector.h>
#include <gsl/gsl_matrix.h>
#include <gsl/gsl_blas.h>
#include <gsl/gsl_multifit_nlinear.h>

/* parameters to model */
struct model_params
{
    double a1;
    double a2;
    double a3;
    double a4;
    double a5;
};

/* Branin function */
int
func_f (const gsl_vector * x, void *params, gsl_vector * f)
{
    struct model_params *par = (struct model_params *) params;
    double x1 = gsl_vector_get(x, 0);
    double x2 = gsl_vector_get(x, 1);
    double f1 = x2 + par->a1 * x1 * x1 + par->a2 * x1 + par->a3;
    double f2 = sqrt(par->a4) * sqrt(1.0 + (1.0 - par->a5) * cos(x1));

    gsl_vector_set(f, 0, f1);
    gsl_vector_set(f, 1, f2);

    return GSL_SUCCESS;
}

int
func_df (const gsl_vector * x, void *params, gsl_matrix * J)
{
    struct model_params *par = (struct model_params *) params;
    double x1 = gsl_vector_get(x, 0);
    double f2 = sqrt(par->a4) * sqrt(1.0 + (1.0 - par->a5) * cos(x1));

    gsl_matrix_set(J, 0, 0, 2.0 * par->a1 * x1 + par->a2);
    gsl_matrix_set(J, 0, 1, 1.0);

    gsl_matrix_set(J, 1, 0, -0.5 * par->a4 / f2 * (1.0 - par->a5) * sin(x1));
    gsl_matrix_set(J, 1, 1, 1, 0.0);

    return GSL_SUCCESS;
}

int
func_fvv (const gsl_vector * x, const gsl_vector * v,
          void *params, gsl_vector * fvv)
{
    struct model_params *par = (struct model_params *) params;
    double x1 = gsl_vector_get(x, 0);
    double v1 = gsl_vector_get(v, 0);
    double c = cos(x1);
    double s = sin(x1);
    double f2 = sqrt(par->a4) * sqrt(1.0 + (1.0 - par->a5) * c);
    double t = 0.5 * par->a4 * (1.0 - par->a5) / f2;

    gsl_vector_set(fvv, 0, 2.0 * par->a1 * v1 * v1);
    gsl_vector_set(fvv, 1, -t * (c + s*s/f2) * v1 * v1);

    return GSL_SUCCESS;
}

void
callback(const size_t iter, void *params,
         const gsl_multifit_nlinear_workspace *w)
{
    gsl_vector * x = gsl_multifit_nlinear_position(w);
    double x1 = gsl_vector_get(x, 0);
    double x2 = gsl_vector_get(x, 1);

    /* print out current location */
    printf("%f %f\n", x1, x2);
}

```

```

void
solve_system(gsl_vector *x0, gsl_multifit_nlinear_fdf *fdf,
             gsl_multifit_nlinear_parameters *params)
{
    const gsl_multifit_nlinear_type *T = gsl_multifit_nlinear_trust;
    const size_t max_iter = 200;
    const double xtol = 1.0e-8;
    const double gtol = 1.0e-8;
    const double ftol = 1.0e-8;
    const size_t n = fdf->n;
    const size_t p = fdf->p;
    gsl_multifit_nlinear_workspace *work =
        gsl_multifit_nlinear_alloc(T, params, n, p);
    gsl_vector * f = gsl_multifit_nlinear_residual(work);
    gsl_vector * x = gsl_multifit_nlinear_position(work);
    int info;
    double chisq0, chisq, rcond;

    printf("# %s/%s\n",
           gsl_multifit_nlinear_name(work),
           gsl_multifit_nlinear_trs_name(work));

    /* initialize solver */
    gsl_multifit_nlinear_init(x0, fdf, work);

    /* store initial cost */
    gsl_blas_ddot(f, f, &chisq0);

    /* iterate until convergence */
    gsl_multifit_nlinear_driver(max_iter, xtol, gtol, ftol,
                               callback, NULL, &info, work);

    /* store final cost */
    gsl_blas_ddot(f, f, &chisq);

    /* store cond(J(x)) */
    gsl_multifit_nlinear_rcond(&rcond, work);

    /* print summary */
    fprintf(stderr, "%-25s %-6zu %-5zu %-5zu %-13.4e %-12.4e %-13.4e (%.2e, %.2e)\n",
           gsl_multifit_nlinear_trs_name(work),
           gsl_multifit_nlinear_niter(work),
           fdf->nevalf,
           fdf->nevaldf,
           chisq0,
           chisq,
           1.0 / rcond,
           gsl_vector_get(x, 0),
           gsl_vector_get(x, 1));

    printf("\n\n");

    gsl_multifit_nlinear_free(work);
}

int
main (void)
{
    const size_t n = 2;
    const size_t p = 2;
    gsl_vector *f = gsl_vector_alloc(n);
    gsl_vector *x = gsl_vector_alloc(p);
    gsl_multifit_nlinear_fdf fdf;
    gsl_multifit_nlinear_parameters fdf_params =
        gsl_multifit_nlinear_default_parameters();
    struct model_params params;

    params.a1 = -5.1 / (4.0 * M_PI * M_PI);
    params.a2 = 5.0 / M_PI;
    params.a3 = -6.0;
    params.a4 = 10.0;
    params.a5 = 1.0 / (8.0 * M_PI);

    /* print map of Phi(x1, x2) */
    {
        double x1, x2, chisq;

        for (x1 = -5.0; x1 < 15.0; x1 += 0.1)

```

```

{
    for (x2 = -5.0; x2 < 15.0; x2 += 0.1)
    {
        gsl_vector_set(x, 0, x1);
        gsl_vector_set(x, 1, x2);
        func_f(x, &params, f);

        gsl_blas_ddot(f, f, &chisq);

        printf("%f %f %f\n", x1, x2, chisq);
    }
    printf("\n");
}
printf("\n\n");
}

/* define function to be minimized */
fdf.f = func_f;
fdf.df = func_df;
fdf.fvv = func_fvv;
fdf.n = n;
fdf.p = p;
fdf.params = &params;

/* starting point */
gsl_vector_set(x, 0, 6.0);
gsl_vector_set(x, 1, 14.5);

fprintf(stderr, "%-25s %-6s %-5s %-5s %-13s %-12s %-13s %-15s\n",
    "Method", "NITER", "NFEV", "NJEV", "Initial Cost",
    "Final cost", "Final cond(J)", "Final x");

fdf_params.trs = gsl_multifit_nlinear_trs_lm;
solve_system(x, &fdf, &fdf_params);

fdf_params.trs = gsl_multifit_nlinear_trs_lmaccel;
solve_system(x, &fdf, &fdf_params);

fdf_params.trs = gsl_multifit_nlinear_trs_dogleg;
solve_system(x, &fdf, &fdf_params);

fdf_params.trs = gsl_multifit_nlinear_trs_ddogleg;
solve_system(x, &fdf, &fdf_params);

fdf_params.trs = gsl_multifit_nlinear_trs_subspace2D;
solve_system(x, &fdf, &fdf_params);

gsl_vector_free(f);
gsl_vector_free(x);

return 0;
}

```

Large Nonlinear Least Squares Example

The following program illustrates the large nonlinear least squares solvers on a system with significant sparse structure in the Jacobian. The cost function is

$$\begin{aligned}
 \Phi(x) &= \frac{1}{2} \sum_{i=1}^{p+1} f_i^2 \\
 f_i &= \sqrt{\alpha}(x_i - 1), \quad 1 \leq i \leq p \\
 f_{p+1} &= \|x\|^2 - \frac{1}{4}
 \end{aligned}$$

with $\alpha = 10^{-5}$. The residual f_{p+1} imposes a constraint on the p parameters x , to ensure that $\|x\|^2 \approx \frac{1}{4}$. The $(p+1)$ -by- p Jacobian for this system is

$$J(x) = \begin{pmatrix} \sqrt{\alpha} I_p \\ 2x^T \end{pmatrix}$$

and the normal equations matrix is

$$J^T J = \alpha I_p + 4xx^T$$

Finally, the second directional derivative of f for the geodesic acceleration method is

$$f_{vv} = D_v^2 f = \begin{pmatrix} 0 \\ 2||v||^2 \end{pmatrix}$$

Since the upper p -by- p block of J is diagonal, this sparse structure should be exploited in the nonlinear solver. For comparison, the following program solves the system for $p = 2000$ using the dense direct Cholesky solver based on the normal equations matrix $J^T J$, as well as the iterative Steihaug-Toint solver, based on sparse matrix-vector products Ju and $J^T u$. The program output is shown below:

Method	NITER	NFEV	NJUEV	NJTJEV	NAEV	Init Cost	Final cost	cond(J)	Final
$ x ^2$ Time (s)									
levenberg-marquardt	25	31	26	26	0	7.1218e+18	1.9555e-02	447.50	2.5044e-01
levenberg-marquardt+accel	22	23	45	23	22	7.1218e+18	1.9555e-02	447.64	2.5044e-01
dogleg	37	87	36	36	0	7.1218e+18	1.9555e-02	447.59	2.5044e-01
double-dogleg	35	88	34	34	0	7.1218e+18	1.9555e-02	447.62	2.5044e-01
2D-subspace	37	88	36	36	0	7.1218e+18	1.9555e-02	447.71	2.5044e-01
stehaug-toint	35	88	345	0	0	7.1218e+18	1.9555e-02	inf	2.5044e-01

The first five rows use methods based on factoring the dense $J^T J$ matrix while the last row uses the iterative Steihaug-Toint method. While the number of Jacobian matrix-vector products (NJUEV) is less for the dense methods, the added time to construct and factor the $J^T J$ matrix (NJTJEV) results in a much larger runtime than the iterative method (see last column).

The program is given below.

```

#include <stdlib.h>
#include <stdio.h>
#include <sys/time.h>
#include <gsl/gsl_vector.h>
#include <gsl/gsl_matrix.h>
#include <gsl/gsl_blas.h>
#include <gsl/gsl_multilarge_nlinear.h>
#include <gsl/gsl_splblas.h>
#include <gsl/gsl_spmatrix.h>

/* parameters for functions */
struct model_params
{
    double alpha;
    gsl_spmatrix *J;
};

/* penalty function */
int
penalty_f (const gsl_vector * x, void *params, gsl_vector * f)
{
    struct model_params *par = (struct model_params *) params;
    const double sqrt_alpha = sqrt(par->alpha);
    const size_t p = x->size;
    size_t i;
    double sum = 0.0;

    for (i = 0; i < p; ++i)
    {
        double xi = gsl_vector_get(x, i);

        gsl_vector_set(f, i, sqrt_alpha*(xi - 1.0));

        sum += xi * xi;
    }

    gsl_vector_set(f, p, sum - 0.25);

    return GSL_SUCCESS;
}

int
penalty_df (CBLAS_TRANSPOSE_t TransJ, const gsl_vector * x,
            const gsl_vector * u, void * params, gsl_vector * v,
            gsl_matrix * JTJ)
{
    struct model_params *par = (struct model_params *) params;
    const size_t p = x->size;
    size_t j;

    /* store 2*x in last row of J */
    for (j = 0; j < p; ++j)
    {
        double xj = gsl_vector_get(x, j);
        gsl_spmatrix_set(par->J, p, j, 2.0 * xj);
    }

    /* compute v = op(J) u */
    if (v)
        gsl_splblas_dgemv(TransJ, 1.0, par->J, u, 0.0, v);

    if (JTJ)
    {
        gsl_vector_view diag = gsl_matrix_diagonal(JTJ);

        /* compute J^T J = [ alpha*I_p + 4 x x^T ] */
        gsl_matrix_set_zero(JTJ);

        /* store 4 x x^T in lower half of JTJ */
        gsl_blas_dsyr(CblasLower, 4.0, x, JTJ);

        /* add alpha to diag(JTJ) */
        gsl_vector_add_constant(&diag.vector, par->alpha);
    }

    return GSL_SUCCESS;
}

```

```

int
penalty_fvv (const gsl_vector * x, const gsl_vector * v,
             void *params, gsl_vector * fvv)
{
    const size_t p = x->size;
    double normv = gsl_blas_dnorm2(v);

    gsl_vector_set_zero(fvv);
    gsl_vector_set(fvv, p, 2.0 * normv * normv);

    (void)params; /* avoid unused parameter warning */

    return GSL_SUCCESS;
}

void
solve_system(const gsl_vector *x0, gsl_multilarge_nlinear_fdf *fdf,
             gsl_multilarge_nlinear_parameters *params)
{
    const gsl_multilarge_nlinear_type *T = gsl_multilarge_nlinear_trust;
    const size_t max_iter = 200;
    const double xtol = 1.0e-8;
    const double gtol = 1.0e-8;
    const double ftol = 1.0e-8;
    const size_t n = fdf->n;
    const size_t p = fdf->p;
    gsl_multilarge_nlinear_workspace *work =
        gsl_multilarge_nlinear_alloc(T, params, n, p);
    gsl_vector * f = gsl_multilarge_nlinear_residual(work);
    gsl_vector * x = gsl_multilarge_nlinear_position(work);
    int info;
    double chisq0, chisq, rcond, xsq;
    struct timeval tv0, tv1;

    gettimeofday(&tv0, NULL);

    /* initialize solver */
    gsl_multilarge_nlinear_init(x0, fdf, work);

    /* store initial cost */
    gsl_blas_ddot(f, f, &chisq0);

    /* iterate until convergence */
    gsl_multilarge_nlinear_driver(max_iter, xtol, gtol, ftol,
                                NULL, NULL, &info, work);

    gettimeofday(&tv1, NULL);

    /* store final cost */
    gsl_blas_ddot(f, f, &chisq);

    /* compute final ||x||^2 */
    gsl_blas_ddot(x, x, &xsq);

    /* store cond(J(x)) */
    gsl_multilarge_nlinear_rcond(&rcond, work);

    /* print summary */
    fprintf(stderr, "%-25s %-5zu %-4zu %-5zu %-6zu %-4zu %-10.4e %-10.4e %-7.2f %-11.4e\n",
            "%-25s %-5zu %-4zu %-5zu %-6zu %-4zu %-10.4e %-10.4e %-7.2f %-11.4e\n",
            gsl_multilarge_nlinear_trs_name(work),
            gsl_multilarge_nlinear_niter(work),
            fdf->nevalf,
            fdf->nevaldfu,
            fdf->nevaldf2,
            fdf->nevalfvv,
            chisq0,
            chisq,
            1.0 / rcond,
            xsq,
            (tv1.tv_sec - tv0.tv_sec) + 1.0e-6 * (tv1.tv_usec - tv0.tv_usec));

    gsl_multilarge_nlinear_free(work);
}

int
main (void)
{
    const size_t p = 2000;

```

```

const size_t n = p + 1;
gsl_vector *f = gsl_vector_alloc(n);
gsl_vector *x = gsl_vector_alloc(p);

/* allocate sparse Jacobian matrix with 2*p non-zero elements in triplet format */
gsl_spmatrix *J = gsl_spmatrix_alloc_nzmax(n, p, 2 * p, GSL_SPMATRIX_TRIPLET);

gsl_multilarge_nlinear_fdf fdf;
gsl_multilarge_nlinear_parameters fdf_params =
    gsl_multilarge_nlinear_default_parameters();
struct model_params params;
size_t i;

params.alpha = 1.0e-5;
params.J = J;

/* define function to be minimized */
fdf.f = penalty_f;
fdf.df = penalty_df;
fdf.fvv = penalty_fvv;
fdf.n = n;
fdf.p = p;
fdf.params = &params;

for (i = 0; i < p; ++i)
{
    /* starting point */
    gsl_vector_set(x, i, i + 1.0);

    /* store sqrt(alpha)*I_p in upper p-by-p block of J */
    gsl_spmatrix_set(J, i, i, sqrt(params.alpha));
}

fprintf(stderr, "%-25s %-4s %-4s %-5s %-6s %-4s %-10s %-10s %-7s %-11s %-10s\n",
    "Method", "NITER", "NFEV", "NJUEV", "NJTJEV", "NAEV", "Init Cost",
    "Final cost", "cond(J)", "Final |x|^2", "Time (s)");

fdf_params.scale = gsl_multilarge_nlinear_scale_levenberg;

fdf_params.trs = gsl_multilarge_nlinear_trs_lm;
solve_system(x, &fdf, &fdf_params);

fdf_params.trs = gsl_multilarge_nlinear_trs_lmaccel;
solve_system(x, &fdf, &fdf_params);

fdf_params.trs = gsl_multilarge_nlinear_trs_dogleg;
solve_system(x, &fdf, &fdf_params);

fdf_params.trs = gsl_multilarge_nlinear_trs_ddogleg;
solve_system(x, &fdf, &fdf_params);

fdf_params.trs = gsl_multilarge_nlinear_trs_subspace2D;
solve_system(x, &fdf, &fdf_params);

fdf_params.trs = gsl_multilarge_nlinear_trs_cgst;
solve_system(x, &fdf, &fdf_params);

gsl_vector_free(f);
gsl_vector_free(x);
gsl_spmatrix_free(J);

return 0;
}

```

References and Further Reading

The following publications are relevant to the **algorithms** described in this section,

- J.J. Moré, *The Levenberg-Marquardt Algorithm: Implementation and Theory*, Lecture Notes in Mathematics, v630 (1978), ed G. Watson.
- H. B. Nielsen, “Damping Parameter in **Marquardt’s** Method”, IMM Department of Mathematical Modeling, DTU, Tech. Report IMM-REP-1999-05 (1999).
- K. Madsen and H. B. Nielsen, “Introduction to Optimization and Data Fitting”, IMM Department of Mathematical Modeling, DTU, 2010.

- J. E. Dennis and R. B. Schnabel, Numerical Methods for Unconstrained Optimization and Nonlinear Equations, SIAM, 1996.
- M. K. Transtrum, B. B. Machta, and J. P. Sethna, Geometry of nonlinear least squares with applications to sloppy models and optimization, Phys. Rev. E 83, 036701, 2011.
- M. K. Transtrum and J. P. Sethna, Improvements to the **Levenberg-Marquardt algorithm** for nonlinear least-squares minimization, arXiv:1201.5885, 2012.
- J.J. Moré, B.S. Garbow, K.E. Hillstom, "Testing Unconstrained Optimization Software", ACM Transactions on Mathematical Software, Vol 7, No 1 (1981), p 17–41.
- H. B. Nielsen, "UCTP Test Problems for Unconstrained Optimization", IMM Department of Mathematical Modeling, DTU, Tech. Report IMM-REP-2000-17 (2000).